



Executive Business Review

Northstar Retail Group — Atlas

Month one in production, and the plan for four more brands

Key messages

- Month one delivered: Tier-1 automation reached 70.1% in the final week, with faster, better-rated answers.
- Quality improved every week, and the evidence is in your dashboard.
- 99.85% availability today: we find the cause in telemetry before buying capacity.
- One base model, one adapter per brand: each brand keeps its voice and its data.
- We're asking for a group sponsor today and Ridgeline as the pilot.

Month one: Tier-1 automation reached 70% by the final week

August averages. The change line compares the first week with the final week, the same basis as your dashboard.

Tier-1 automation 67.9% +4.3 pts since week 1 Final week 70.1%; 70% reached from Aug 27	Avg handle time 7m 44s 8m 03s → 7m 25s, week 1 → week 4 –35% vs pre-launch baseline, supplied by Northstar	CSAT 82.3 80.8 → 83.7, week 1 → week 4 +12 pts vs pre-launch baseline, supplied by Northstar	Requests / day 25.5K +45% since week 1 789,563 in August
Availability 99.85% 99.9% met on 4 of 31 days Lowest 99.69% on Aug 9	P50 / P95 latency 605 ms / 1.66 s P95 –368 ms since week 1 P50 –134 ms since week 1	Eval pass rate 93.1% +3.3 pts since week 1 Grounded answers 93.9%	Escalation rate 12.8% –2.7 pts since week 1 Final week 11.5%

The –35% handle-time and +12-point CSAT figures compare against Northstar’s pre-launch baselines, which are not in the August dataset; the in-month trends are measured directly.

Operational health is 3.58 / 5, led by quality and outcomes

Four operational pillars, each scored 1–5 on full-period averages for August. The same scoring as your dashboard.

Operational health, 1–5

3.58 / 5

Basis

Full-period averages (Aug 1 – Aug 31, 31 days)



Reliability

Is the platform up, error-free and fast?



Adoption

Is Northstar putting more of its support load on the model?



Model quality

Are the answers right?



User outcomes

What did Northstar's shoppers experience?




Quality improved every week, and adoption is still rising

Daily values for August, with the operational events marked as on your dashboard.


Tier-1 automation rate

Rose from 65.8% in week 1 to 70.1% in the final week, while automated tickets grew +15% against total tickets +8%.



Grounding and eval pass

Grounded answers 92.2% → 96.0%, eval pass 91.4% → 94.8%, escalations -2.7 pts (week 1 → week 4).



Worth reviewing together


Requests per Tier-1 ticket rose from 4.6 to 6.2 over the month (7-day averages). It may be longer conversations, repeat contacts or non-ticket traffic; we'd like to look at it with your team.

Find the cause before buying capacity

Availability averaged 99.85% and met 99.9% on 4 of 31 days, so the gap is not only the two incident days. Daily metrics can't show why; deployment telemetry can.

Availability

Averaged 99.85%; 99.9% met on 4 of 31 days, all in the final week. 99.9% is a proposed target; no SLA was in place for August.



Recommended now

Diagnose and tune

Review deployment telemetry together (scale-up events, error codes by hour), retune scale-up thresholds and pre-warm before promotions.
Bill stays approx. \$1.4k/month at current traffic

Only if the data shows it

Always-on capacity for Atlas alone

Approx. \$5,800/month, or \$8,800 with EU-only placement.
Justified only if scale-up is driving the misses

With expansion


Shared five-brand deployment

Always-on capacity shared by 5 brands, billed as dedicated capacity: approx. \$1,200–\$1,800 per brand per month.
Baseline capacity is affordable once shared

Capacity figures at Fireworks list rates as of 5 Oct 2026; contracted terms may differ.

One base model, one adapter per brand, data kept apart

Each brand keeps its own voice, knowledge and data on a single shared deployment.



Data isolation

Each brand's catalogue and policy facts sit in its own knowledge index; each adapter is trained on that brand's transcripts.

Why not reuse Atlas with a new prompt?

Atlas's tone and policy-handling habits are learned into its adapter and would leak into other brands.

Why not five separate models?

Five deployments to run and retrain, and the smaller brands lack the data to train them well.

Ridgeline pilot: two arms, criteria agreed up front

Closest to Atlas in product and customer, on the same Zendesk tooling, at moderate volume: it tests tone and return-policy handling with limited exposure.

Arm A · Lowest cost

Base model + Ridgeline system prompt + Ridgeline knowledge index

No brand-specific training. If it clears every criterion, no adapter is needed.

Arm B · Brand-tuned

Base model + Ridgeline adapter + Ridgeline system prompt + knowledge index

Tone, workflow and policy-following learned from Ridgeline transcripts; policy facts stay in the knowledge index.

Success criterion, four weeks

Target

Grounded answers

≥ 95%

Eval pass rate

≥ 93%

Escalation rate

≤ 12%

CSAT

≥ Atlas final week (83.7)


P95 latency

≤ 1.3 s, held through one promotion peak

Wave 2 goes ahead when one arm clears every criterion. The arm is chosen on quality, and Ridgeline's handle-time and CSAT baselines are captured before launch so the uplift is measured directly.

From pilot to five brands by mid-2027

Each step has a gate; wave 2 starts only after the pilot readout.



Capacity figures at Fireworks list rates as of 5 Oct 2026; contracted terms may differ.

Biggest risk: the rollout stalls at one brand

The risk is commercial and organisational rather than technical: month one worked, but no one owns the rollout above brand level yet.

Biggest risk

Without group-level sponsorship the rollout stays with Atlas, and shared capacity never carries enough volume to make per-brand cost efficient.

Why it matters


- Always-on capacity stays a single-brand cost (approx. \$5,800/month) rather than \$1,200–\$1,800 per brand shared.
- The four sister brands, with most of the group's Tier-1 volume, don't get the gains Atlas has shown.
- Each brand would make its own platform decision, so the group ends up with several tools to run.

How we mitigate it

- **Name a group sponsor at this review** 22 Oct
Someone above brand level who owns the rollout and the wave-2 decision.
- **Agree the Ridgeline pilot criteria together** 22 Oct
Success is defined jointly up front, so the readout is a decision rather than a debate.
- **Size shared capacity with a second brand** Nov 2026
Capacity is committed with Ridgeline from a load benchmark, never for Atlas alone.

Thank you

Questions and discussion



Joint next actions

Action	Owner	Due
Retune scale-up thresholds and agree a pre-warm runbook for planned promotions	Fireworks Engineering + Northstar CX operations	16 Oct
Share pre-launch handle-time and CSAT baselines for Atlas	Deployment Strategist + Northstar CX	16 Oct
Executive review: agree a group sponsor, the Ridgeline pilot and its success criteria	Deployment Strategist + Account Executive	22 Oct
Joint review of rising requests per ticket (4.6 → 6.2)	Deployment Strategist + Northstar Engineering	23 Oct
Train the Ridgeline adapter and run the two-arm pilot	Fireworks Applied ML + Northstar Engineering	18 Dec
Pilot readout and wave-2 decision	Deployment Strategist + Northstar sponsor	8 Jan 2027