

Atlas automated 70% of Tier-1 support by month end



70.1%

Tier-1 automation, final week

67.9% August average;
70% reached from 27 Aug



7.41 min

Avg handle time, final week

8.05 min in week 1; –35%
vs Northstar's pre-launch
baseline



83.7

CSAT, final week

80.8 in week 1; +12 pts vs
Northstar's pre-launch
baseline



99.85%


Availability, August

Dips on 9 Aug (catalogue
sync) and 21 Aug
(promotion traffic)

In-month figures are measured from August data (week 1 = 1–7 Aug; final week = 25–31 Aug). The –35% and +12-point figures compare against Northstar's pre-launch baselines, which are not in the August dataset.

Quality improved every week; adoption is still rising

Answer quality by week



■ Adoption is still rising

Automation 65.8% → 70.1% from week 1 to the final week, with no plateau yet.

■ Fewer handoffs to agents

Escalations 14.2% → 11.5% as grounding improved after the 24 Aug knowledge refresh.

■ Availability missed 99.9% most days

Met on 4 of 31 days, all in the final week. The fix has a cost (next slide).

■ Requests per ticket 4.6 → 6.2

Worth a joint review, not an alarm: longer conversations, repeat contacts or non-ticket traffic.

99.9% arrives with shared capacity, not a 5x bill

Atlas scales its capacity down when traffic is quiet. That keeps the bill low, and it explains most of August's availability misses: after a quiet spell, the first requests wait while capacity comes back up.

■ RECOMMENDED NOW

Tune what Atlas has

- Retune scale-up thresholds
- Pre-warm capacity before planned promotions

Bill stays ~\$1.4k/month at current traffic; availability ~99.85%

■ NOT RECOMMENDED

Always-on capacity for Atlas alone

- ~\$5.8k/month
- ~\$8.8k/month with EU-only placement

Buys about 22 minutes a month of extra availability

■ WITH EXPANSION


Shared five-brand deployment

- Always-on capacity shared by five brands
- ~\$1.2k–1.8k per brand per month

This is where 99.9% becomes efficient

Capacity figures at Fireworks list rates as of 5 Oct 2026; contracted terms may differ.

One base, one adapter per brand, data kept apart



■ Data isolation

Each brand's catalogue sits in its own knowledge index; each adapter is trained on that brand's transcripts.

■ Why not reuse Atlas with a new prompt?

Atlas's tone and return policies are learned into its adapter and would leak into other brands.

■ Why not five separate models?

Five deployments to run and retrain, and the smaller brands lack the data to train them well.

Ridgeline pilot: two arms, agreed criteria

Ridgeline is closest to Atlas in product and customer, uses the same Zendesk tooling and runs at moderate volume, so it tests tone and return-policy transfer with limited exposure.

■ ARM A · LOWEST COST

Base model + Ridgeline prompt + Ridgeline knowledge index

No brand-specific training. If it clears every criterion, no adapter is needed.

■ ARM B · BRAND-TUNED


Base model + Ridgeline adapter + prompt + knowledge index

Tone and policy learned from Ridgeline transcripts.

Success criterion (four weeks)	Target
Grounded answers	≥ 95%
Eval pass rate	≥ 93%
Escalation rate	≤ 12%
CSAT	≥ Atlas final week (83.7)
P95 latency	≤ 1.3 s through one promotion peak

Evidence for wave 2: one arm clears every criterion, the arm is chosen on quality, and Ridgeline’s handle-time and CSAT baselines are captured before launch so the uplift is measured directly.

From pilot to five brands by mid-2027



The rollout must not stall at one brand

■ BIGGEST RISK

Without group-level sponsorship the rollout stays with Atlas, and shared capacity never carries enough volume to make 99.9% and per-brand cost efficient.

MITIGATION

- Name a group sponsor at this review
- Agree the Ridgeline pilot criteria together
- Commit always-on capacity alongside a second brand

KEY MESSAGES

- 01** Month one delivered: 70% of Tier-1 automated by month end, faster and better-rated.
- 02** Quality improved every week, and the evidence is in your dashboard.
- 03** 99.85% is what Atlas alone buys today; 99.9% comes with shared capacity.
- 04** One base, one adapter per brand: each brand keeps its voice and its data.
- 05** We're asking for a group sponsor today and Ridgeline as the pilot.