

Northstar dashboards: design document

Northstar Retail Group · Atlas customer-support deployment · August 2026 data

Two dashboards are built from the same 31-day dataset. The **internal QBR** is for Fireworks' account, engineering and leadership teams. The **customer account-health dashboard** is for Northstar's VP Customer Experience and VP Engineering. They are two separate apps sharing one data, scoring and design package, so internal judgement cannot reach the customer build. A test enforces this.


The sections follow the brief:

1. Internal QBR screenshots
 2. Customer dashboard screenshots
 3. Audience and purpose of each view
 4. Why these metrics and visualisations
 5. How the account-health score is calculated
 6. How internal and customer content differ
 7. What each view intentionally excludes
 8. Assumptions, caveats and design trade-offs
-


1. Internal QBR dashboard

Three views behind a menu next to the logo: **Account health**, **Performance trends** and **Expansion plan**. A **Single page** switch stacks all three for reading end to end.


Account health. The customer score, its six pillars and the eight headline metrics. Each tile opens its chart.



Focus areas. Opened from the score panel: the three items furthest from a perfect 5, with what it takes to reach the next band.



Performance trends. Nine daily charts. Every chart marks the three operational events (9, 21 and 24 Aug); red marks incidents and purple marks changes.



Expansion plan. The pipeline across the four sister brands, followed by risks, stakeholder coverage, dependencies and margin, competitive risk, decisions needed and next actions.

Expansion pipeline

Brands live 1 of 5 Atlas only	Tier-1 volume in pipeline 7.0× Atlas 32,600 tickets/day across 4 brands	Shared deployment value \$70,000+/yr One H100 at list (\$105,000 EU), dedicated billing; vs \$14,000/yr per token today	Pilot Ridgeline Nov–Dec 2026
---	---	---	--

Brand	Stage	Tier-1 / day	Catalogue	Tone · returns	Tooling	Likelihood on timeline	Rationale
Atlas Flagship. Premium technical outerwear	Live Live since Aug 2026	4,627 from data	~3k SKUs, spec-heavy	Expert, understated · 30-day returns	Zendesk	—	Adapter loads unchanged onto the shared deployment
Ridgeline Mountain and technical outerwear	Pilot Nov–Dec 2026	6,200 1.3× Atlas	~4k SKUs, spec-heavy	Adventurous · 60-day returns	Zendesk	High	Closest to Atlas, same tooling, moderate volume: limits blast radius while testing tone and policy handling
Halden Mass-market everyday outerwear	Wave 2 Q1 2027	13,900 3.0× Atlas	~18k SKUs, broad	Friendly, direct · 30-day returns	Salesforce Service Cloud	Medium	Largest volume and value; new tooling integration; Agentforce is in the account
Polar Children's outerwear	Wave 2 Q1 2027	7,100 1.5× Atlas	~6k SKUs, size-led	Warm, reassuring · 90-day returns	Gorgias	Medium	Sizing and product-safety queries need strict grounding
Harbour & Hide Heritage waxed and leather, repairs	Wave 3 Q2 2027	5,400 1.2× Atlas	~2k SKUs plus repair services	Formal, heritage · repairs and warranty	In-house	Low	Repair workflows and in-house tooling; prompt + RAG arm may suffice

Sister-brand profiles are scenario assumptions. Shared deployment value is one dedicated H100 at list, billed to Northstar per GPU-second after migration; the replica count for five brands' combined load has to be benchmarked, so it is a floor, not a forecast.

Top three risks

Type	Risk	Evidence	Mitigation	Owner
Technical	Availability below 99.9% on most days; cause not yet established	27 of 31 days below 99.9%, 25 of them outside the 9 and 21 Aug events; P95 spikes on those two days point to burst and scale-up behaviour there, but daily metrics cannot separate scale-up, application errors and other serving failures	Review deployment telemetry before buying capacity; retune scale-up thresholds and pre-warm for known promotions now; add a warm floor only if telemetry shows cold starts or scale-up lag are a material cause (see Decisions)	Fireworks Engineering
Relationship	No sponsor above brand level, so the account stays single-brand	Group executive and all four sister-brand CX leads are gaps in stakeholder coverage; Partnership scores 2	Use the EBR to secure a named group sponsor and agree the Ridgeline pilot and success criteria jointly	Deployment Strategist + Account Executive
Execution	Pilot slips on data: Ridgeline transcripts, catalogue export and legal sign-off on de-branded Atlas data	Three Northstar-owned dependencies due 6 Nov, none started; without them arm B cannot train and arm A cannot ground	Dependencies tabled at the EBR with owners; fallback is arm A (prompt + RAG, no adapter) if transcripts are late	Deployment Strategist + Northstar CX operations

One technical, one relationship, one execution risk, as the brief asks. The "Focus areas" pop-up on the health view is the scoring-derived list and is a different thing.

Stakeholder coverage

Stakeholder	Status	Notes
Atlas CX operations lead	Champion	Day-to-day owner; drove the move to production
VP Customer Experience (group)	Engaged	EBR audience; candidate business sponsor for the rollout
VP Engineering (group)	Engaged	EBR audience; owns routing, RAG and integration work
Group executive (COO / CDO)	Gap	Needed for group-level approval and budget
Sister-brand CX leads (x4)	Gap	No contact yet; pilot brand first
Legal / data protection	Gap	Sign-off for de-branded Atlas transcripts

Decisions needed

01	Fund the Ridgeline pilot Recommendation Yes: Applied ML time for one adapter and eval set at no charge, in exchange for a named group sponsor and agreed success criteria	Field CTO, Europe By 16 Oct
02	Commercial structure for the group rollout Recommendation Atlas stays on per-token billing until it migrates. The shared multi-LoRA deployment is billed as dedicated capacity (per GPU-second) from the pilot, with committed spend from wave 2; size replicas from a load benchmark before committing	Account Executive + Sales leadership By 20 Oct
03	Reserve BF16 capacity before commitment Recommendation Reserve for the pilot window only (weeks, not months); release if the pilot slips	Fireworks Infrastructure lead By 30 Oct
04	Availability target: Atlas now vs the shared deployment Recommendation Diagnose before buying capacity. The gap to 99.9% is spread over 27 of 31 days and its cause is not established; at 19–31k requests a day Atlas is unlikely to sit at zero for long. Pull deployment telemetry, retune autoscaling and pre-warm for known promotions; add a warm minimum replica (~\$4.6k/month extra, \$7.6k with EU placement) only if telemetry shows cold starts or scale-up lag materially drive the misses. On the shared deployment a floor costs ~\$1.2k–1.8k per brand	Field CTO + Engineering By 20 Oct


Next actions

Action	Owner	Due	Status
Pull Atlas's deployment telemetry (replica count, scale-up events, 503s and error codes by hour), confirm its serving set-up and how reported spend maps to the bill	Deployment Strategist + Fireworks Engineering	9 Oct	In progress
Retune scale-up thresholds and write the promotion pre-warm runbook (warm floor only if telemetry calls for it)	Fireworks Engineering + Northstar CX operations	16 Oct	Not started
Obtain pre-launch handle-time and CSAT baselines to substantiate the ~35% and +12 claims	Deployment Strategist + Northstar CX	16 Oct	Not started
Hold the EBR: secure a group sponsor, agree the Ridgeline pilot and success criteria	Deployment Strategist + Account Executive	22 Oct	Not started
Joint review of rising requests per ticket (4.6 → 6.2)	Deployment Strategist + Northstar Engineering	23 Oct	Not started
Stand up the shared BF16 deployment, load the Atlas adapter, confirm parity on the eval set	Fireworks Engineering	13 Nov	Not started
Train the Ridgeline adapter and run the two-arm pilot	Fireworks Applied ML + Northstar Engineering	18 Dec	Not started
Pilot readout and wave-2 go / no-go	Deployment Strategist + Northstar sponsor	8 Jan 2027	Not started


2. Customer account-health dashboard

Five views: Outcomes, Service reliability, Answer quality, Spend and Next steps. It has the same Single page switch.

Outcomes. Operational health on the four measured pillars, then eight headline metrics: the internal set, with Tier-1 automation in place of spend.



Service reliability. Availability against a proposed 99.9% target, error rate and latency, plus a Known caveats card that states what is not yet known.



Known caveats

Availability target

99.9% was met on 4 of 31 days, and the misses were spread across the month rather than limited to the 9 and 21 Aug events. The cause is not yet established: possible contributors include capacity scale-up during bursts, application errors and other serving failures. We will review the deployment telemetry with you before recommending any change in capacity; in the meantime we are retuning scale-up thresholds and pre-warming before planned promotions.

Requests per ticket


Requests per Tier-1 ticket rose from 4.6 to 6.2 over the month (7-day averages). We would like to review this together: it may be longer conversations, repeat contacts or non-ticket traffic.

Headline improvements

The ~35% handle time and +12-point CSAT figures compare against Northstar's pre-launch baselines, which are not in this dataset. The in-month trends shown here are measured directly.

Answer quality. Grounding, evaluation pass rate and escalations; CSAT and handle time.

Answer quality



Spend. Daily and cumulative spend, unit cost, and the budget context for the capacity decisions ahead.

Budget context

August

\$1,213, billed per token (approx. \$1.12 per 1M tokens).

Current run-rate

Approx. \$1,400/month at final-week traffic.

Always-on capacity

Holding a replica warm around the clock would move Atlas to approx. \$5,800/month (\$8,800 with EU-only placement). We would recommend it only if the deployment telemetry shows that capacity scale-up is driving the availability misses.

On the shared deployment

The shared 5-brand deployment is billed as dedicated capacity; one always-on replica works out at approx. \$1,200-\$1,800 per brand per month.

Capacity figures use Fireworks list rates as of 5 Oct 2026; contracted terms may differ.

Next steps. The Ridgeline pilot: why Ridgeline, the two arms and the success criteria. Below it on the page: what Fireworks needs from Northstar, and the joint actions with owners and dates.

Pilot: Ridgeline

Window

Nov-Dec 2026, four weeks of live traffic

Why Ridgeline

Closest to Atlas in product and customer, on the same Zendesk tooling, at moderate volume: it tests tone and return-policy handling with limited exposure.

Arm	Set-up	Purpose
Arm A	Base model + Ridgeline system prompt + Ridgeline knowledge index	Lowest cost: no brand-specific training
Arm B	Base model + Ridgeline adapter + Ridgeline system prompt + knowledge index	Brand-tuned: tone, workflow and policy-following behaviour learned from Ridgeline transcripts; policy facts stay in the knowledge index

Success criterion	Target
Grounded answers	≥ 95%
Eval pass rate	≥ 93%
Escalation rate	≤ 12%
CSAT	≥ Atlas week - 4 (83.7)
P95 latency	≤ 1.3 s, held through one promotion peak

Wave 2 goes ahead when one arm clears every criterion over four weeks. The arm is chosen on quality, and if Arm A is enough no adapter is trained. Ridgeline baselines (handle time, CSAT) are captured before launch so the uplift can be measured directly this time.

3. Audience and purpose of each view

View	Audience	Purpose
Internal: Account health	Account team, engineering, leadership	One number for the account, what is dragging it down, and the headline metrics behind both
Internal: Performance trends	Engineering, account team	The evidence: daily trends with operational events marked, so a spike is always read next to its cause
Internal: Expansion plan	Leadership, account team	Where to focus: pipeline and its value, top three risks, stakeholder gaps, dependencies, margin, competitive threats, decisions needed and actions
Customer: Outcomes	Northstar VP CX, VP Engineering	Demonstrated results: operational health and the headline metrics, each with its caveat
Customer: Service reliability	VP Engineering	Operational transparency, including what we do not yet know about availability and what we are doing about it
Customer: Answer quality	VP CX	Whether answers are right and getting better, and what shoppers experienced
Customer: Spend	VP CX, VP Engineering, finance	The bill in context, and what always-on capacity would cost alone versus shared
Customer: Next steps	Both VPs	The pilot proposal, Northstar's dependencies and the joint plan

The split follows the standard QBR/EBR distinction. A QBR answers "are we on track, and what is blocking us?", so the internal dashboard is a working tool: operational detail, judgement and an action list. An

executive review asks "so what?" and leads with outcomes, so the customer view leads with results, keeps operational detail one click away, and ends on a joint plan.

4. Why these metrics and visualisations

Sources of inspiration

- **The brief's three outcomes:** Tier-1 automation, handle time and CSAT appear on both dashboards. Each carries its caveat in the same place: 70% is the month-end figure (the August average is 67.9%), and the -35% and +12-point improvements depend on Northstar's pre-launch baselines, which are not in the dataset.
- **QBR practice:**
 - usage and health metrics, SLA performance, a health score;
 - open issues with owners, and dated actions on both sides;
 - avoiding the common failure of a "data dump" of every chart.

That shaped the tile-first layout and the plan view.

- **EBR practice:** surface a metric only if it answers "so what?". This is why the customer view leads with outcomes and puts caveats next to each claim.
- **GitLab's customer health scoring (PROVE)** for the structure of the health score (section 5).
- **Published benchmarks** to calibrate what "good" means for each metric (section 5): the Freshworks Customer Service Benchmark 2025 (Retail & eCommerce), Salesforce's CSAT guidance and State of Service 2025, and the evaluation-framework practice of RAGAS and Microsoft Foundry.
- **The dataset itself:** daily figures for one deployment, so every metric is a daily series or an average of one. There is no per-conversation drill-down.

Metric by metric

Metric	Why it is shown	Where
Tier-1 automation	The brief's headline outcome and the clearest measure of adoption: the share of Tier-1 work entrusted to the model	Customer tile, chart, Adoption pillar
Avg handle time	Brief outcome; the efficiency gain for Northstar's agents. Shown as minutes and seconds (7m 44s), because "7.74 min" reads as a decimal	Tile, chart, User outcomes
CSAT	Brief outcome; what shoppers felt	Tile, chart, User outcomes
Requests / day	Scale of use. Shown, but not scored: requests per ticket rose from 4.6 to 6.2, so request volume measures consumption, not adoption	Tile, chart
Availability, error rate	Is the platform up and error-free? The proposed 99.9% target is labelled as proposed; no SLA was in place	Tile, charts, Reliability
P50 / P95 latency	Speed of a chat answer. P50 is the typical case; P95 is the burst behaviour that promotions expose	Tile, chart, Reliability
Grounded answers, eval pass	Are the answers right? Direct measures of the model, kept separate from customer-experience measures	Tile, chart, Model quality

Metric	Why it is shown	Where
Escalation rate	Hand-offs to agents. Scored as first-contact resolution (100% – escalation) because it has a direct published benchmark	Tile, chart, User outcomes
Spend, cost per 1k requests	The bill and unit cost; unit cost held steady as usage grew	Internal tile, charts, customer Spend view

Visualisation choices

- **Headline tiles first, charts one click away.** Each tile shows the August average. A change line underneath compares week 1 with the final week, coloured teal when it improves. Leaders read the tiles; engineers follow the link to the chart.
- **Line charts for rates and latency; bars plus a cumulative line for spend.** A second axis is used only when two related units tell one story (requests vs tickets, CSAT vs handle time).
- **Event markers on every chart,** so the 9 Aug catalogue-sync incident and the 21 Aug promotion burst are never mistaken for trends. Reference lines are dashed and labelled (70% automation, 99.9% proposed target).
- **The Fireworks design language,** taken from fireworks.ai's own CSS: colour tokens, the framed header, Inter type and square status pills. One colour is off-brand: amber, for the Positive-Watch status.

5. How the account-health score is calculated

RAMP UP is adapted from GitLab's customer health scoring framework, **PROVE** (Product, Risk, Outcomes, Voice of the customer, Engagement). It was rebuilt from a Fireworks Deployment Strategist's point of view.

PROVE is written for a SaaS seat-based product. An inference account has questions PROVE doesn't ask:

- Is the serving platform reliable?
- Is the model right?
- Is the account worth what it costs Fireworks to serve?

How PROVE maps to RAMP UP:

- **Product** becomes **Adoption**.
- **Risk** and **Engagement** become **Partnership**, read as readiness to expand beyond one brand.
- **Outcomes** and **Voice of the customer** become **User outcomes**.
- **Reliability**, **Model quality** and **Profitability** are the inference-specific additions.

The pillars are mutually exclusive, and every one is scored so that higher means healthier.

	Pillar	Question it answers	Weight	Basis
R	Reliability	Is the platform up, error-free and fast?	20%	Measured
A	Adoption	Is Northstar putting more of its support load on the model?	20%	Measured
M	Model quality	Are the answers right?	20%	Measured
P	Partnership	How committed is Northstar beyond the first deployment?	10%	Account-team judgement

	Pillar	Question it answers	Weight	Basis
U	User outcomes	What did Northstar's shoppers experience?	20%	Measured
P	Profitability	Is the account commercially worth it to Fireworks?	10%	Account-team judgement

Scoring rules

- Metric scores.** Each metric scores 1–5 against bands. A value on a band edge earns the higher score.
- Full-period averages.** Every metric uses the average of all 31 days, not a trailing week. This is fair to the numbers rather than flattering: scored on the strong final week alone, the account would read 3.73 internally and 4.17 on the customer view, against 3.27 and 3.58 for the month.
- Rounding.** Values are rounded to display precision before scoring, so the number on screen is the number that was scored.
- Pillar scores.** A pillar is the mean of its metric scores. P50 and P95 latency count half each, so latency counts once.
- Overall score and status.** The overall score is the weighted sum of the pillars, on a 1–5 scale. Status uses the same scale: **Healthy** ≥ 4, **Positive-Watch** 3–3.99, **Negative-Watch** 2–2.99, **At risk** below 2.
- Weights.** Measured pillars carry 80% of the weight, so evidence dominates. If a pillar has no input, its weight is shared among the others, as in GitLab's model.

Bands and where they come from

Metric	5	4	3	2	Calibration
Availability	≥ 99.9%	≥ 99.7%	≥ 99.5%	≥ 99.3%	Account-team bands (no SLA was supplied); 99.9% = 5, then 0.2-point steps
Error rate	≤ 0.5%	≤ 1%	≤ 2%	≤ 5%	Observed server-error rates of major LLM APIs (~0.3–0.7%)
P50 latency	≤ 300 ms	≤ 500 ms	≤ 800 ms	≤ 1.2 s	Time-to-first-token guidance: under 500 ms feels instant
P95 latency	≤ 1.0 s	≤ 1.5 s	≤ 2.0 s	≤ 3.0 s	As above, for the burst tail
Tier-1 automation	≥ 80%	≥ 70%	≥ 60%	≥ 45%	Account-team bands, set against Freshworks (53% retail AI deflection) and Salesforce (30% of cases AI-handled today, 50% by 2027)
Grounded answers, eval pass	≥ 95%	≥ 90%	≥ 85%	≥ 80%	Account-team quality gates. RAGAS's example CI gates (0.95 / 0.90) set the 5 and 4 edges. Microsoft Foundry's agent evaluators supply the structure: component checks, plus a composite that passes only if every component passes
CSAT	≥ 85	≥ 70	≥ 60	≥ 50	Salesforce: above 70% is "good", below 50% "less desirable"
Handle time	≤ 2m 03s	≤ 12m 07s	≤ 1h 08m	≤ 2h 00m	Freshworks 2025, Retail & eCommerce conversational resolution time (Trendsetter / Performer / Aspirant)

Metric	5	4	3	2	Calibration
First-contact resolution	≥ 93.95%	≥ 82.43%	≥ 68.12%	≥ 58.00%	Freshworks 2025, same tiers (score 2 is our extension)

No framework publishes a five-step scale for every metric. Where the bands are ours, the table says so, and names the published figure they were anchored to.

August result

Pillar	Inputs → metric scores	Pillar score
Reliability	Availability 99.85% → 4 · error rate 1.04% → 3 · P50 605 ms → 3 · P95 1.66 s → 3	3.33
Adoption	Automation 67.9% → 3 (ended the month at 70.1%)	3.00
Model quality	Grounded 93.9% → 4 · eval pass 93.1% → 4	4.00
Partnership	Live on one brand of five; no sponsor above brand level yet	2
User outcomes	CSAT 82.3 → 4 · handle time 7m 44s → 4 · first-contact resolution 87.2% → 4	4.00
Profitability	Revenue currently immaterial (~\$14k a year); expected to improve as Northstar moves to dedicated capacity on the shared deployment	2
Overall (internal)		3.27 / 5, Positive-Watch
Operational health (customer)	R, A, M and U only, weights redistributed to 25% each	3.58 / 5

Focus areas (internal only) lists the three items furthest from a perfect 5. Items are ranked by their exact position within each band, not by whole-number score: at 605 ms, P50 latency sits 65% of the way from a 3 to a 4. For August the three are Partnership, Profitability and P50 latency. This list is scoring-derived. The brief's top three risks (technical, relationship, execution) are a separate register on the Expansion plan view.

6. How internal and customer content differ

Item	Internal QBR	Customer dashboard
Health score	All six pillars, 3.27, with status band, pillar letters and weights	Four measured pillars, 3.58, "Operational health, 1–5", pillar icons, no weights or status band
Metric notes	Assumptions and benchmark notes under each metric	Hidden: methodology is ours to explain, not theirs to read
Focus areas	Yes	No
Headline metrics and trend charts	Yes; spend among the eight tiles	Yes, the same figures from the same code; Tier-1 automation in place of spend
Operational events	As logged	Restated in plain customer language

Item	Internal QBR	Customer dashboard
Risks	Top three risks register; competitive risk	Known caveats card: what we don't yet know and what we're doing about it
Expansion	Pipeline with likelihood on timeline, stakeholder coverage, dependencies on both sides	Pilot proposal and Northstar's dependencies only
Money	Margin, deployment economics, shared-deployment value	Budget context: the bill, run-rate, and what always-on capacity would cost alone and shared
Plan	Leadership decisions needed; every action	Joint actions only, with the same dates and owners as the internal list (a test keeps them in step)

7. What each view intentionally excludes

From the customer dashboard. The brief bars internal sentiment, competitive intelligence, margin and expansion confidence, so the customer dashboard leaves out:

- the Partnership and Profitability scores;
- the status band and Focus areas;
- likelihood on timeline and stakeholder status;
- competitors, margin and commercial-structure decisions.

A boundary test fails the build if the customer app imports internal code or uses internal-only vocabulary, so this can't regress.

From the internal dashboard: methodology prose. The page carries facts and evidence; explanation lives in this document.

From both dashboards:

- **Named benchmark sources.** They calibrated the bands; on the page they would add clutter, not evidence.
- **Per-conversation drill-down.** The data is daily.
- **Pre-launch baselines.** Not supplied, so the -35% and +12-point figures are always attributed to Northstar.
- **Any SLA claim.** None was in place for August.
- **Animated figures.** Numbers render as fixed facts. Motion is limited to a soft view entrance and hover states, with a reduced-motion fallback.

8. Assumptions, caveats and design trade-offs

Assumptions about the data

- **Billing.** `spend_usd` tracks tokens at a constant ~\$1.12 per 1M on every day, so Atlas is taken to be billed per token today. On the proposed shared deployment, Northstar moves to dedicated capacity billed per GPU-second. GPU utilisation can't be derived from token billing, so none is quoted; it needs deployment telemetry or a load benchmark.
- **Latency.** The latency columns are assumed to be time-to-first-token: 331 output tokens in ~600 ms end to end would be implausible.

- **CSAT and eval pass.** CSAT is assumed to be % satisfied. The eval pass rate is supplied without its rubric, so it is treated as an end-to-end quality gate for planning; the rubric is a listed dependency.
- **Scenario.** Northstar is an outerwear group, and Atlas is the flagship. The four sister brands (Ridgeline, Halden, Polar, Harbour & Hide) and their volumes are invented, because the brief's brand descriptions were not supplied.
- **Pricing.** From fireworks.ai/pricing, checked 5 Oct 2026: H100 \$8.00/GPU-hour and region-restricted deployments 1.5x. Every price lives in one file (`shared/src/pricing.ts`); re-check before presenting.

Caveats shown on the dashboards

- **Automation.** 70% is the month-end level; the August average is 67.9%.
- **Baselines.** Handle-time and CSAT improvements versus baseline use Northstar's figures.
- **Availability.** Availability missed 99.9% on 27 of 31 days, 25 of them outside the two incidents. The cause is not established from daily metrics.
- **Requests per ticket.** Requests per ticket rose from 4.6 to 6.2 and is flagged for a joint review.

Design trade-offs

Choice	Over	Why
Two separate apps sharing one package	One app with role-based views	The customer build cannot contain internal data, so there is no access control to get wrong. Cost: some duplicated view code
Full-period averages for every score and tile	A trailing 7-day window	Fair rather than flattering: the final week alone scores 3.73 internally (4.17 for the customer) versus 3.27 (3.58) for the month. The trend is shown separately as change lines and charts
Adoption scored on automation penetration only	Also scoring request-volume growth	More requests per ticket would have scored as more adoption, the opposite of what matters. Automated-ticket growth (+15% vs total tickets +8%) is shown as context
No rate-of-change score for adoption	Crediting August's steady climb (+5.4 points a month)	Growth has a ceiling: a highly automated account has little room left, so a pace score would penalise the best accounts
Judgement pillars at half weight	Equal weights	Evidence should dominate. The two judgement calls are visible and explained, not hidden in the average
Profitability from Fireworks' side	Customer return on investment	Customer value already shows in User outcomes and the EBR. Scored 2: low today, but with a clear path as Northstar moves to dedicated capacity
Lenient availability bands (99.9% = 5)	Anchoring to a 99.99% enterprise SLA	No SLA was in place. Anchored to 99.99%, August would score a 2. Both scales are documented
CSAT on Salesforce's good/poor guidance	Freshworks' retail CSAT tiers	Freshworks' tiers come from surveys of mostly human-agent chats, too harsh a yardstick for an AI agent's first month
Escalation scored as first-contact resolution, in User outcomes	Scoring it under Model quality	It has a direct published benchmark, and Model quality stays about whether answers are right
Telemetry before capacity	Recommending an always-on replica to reach 99.9%	Daily metrics can't show why availability missed. A warm replica is ~\$5.8k/month for Atlas alone versus ~\$1.2–1.8k per brand when shared, so it is recommended only if telemetry shows cold starts or scale-up are the cause

Choice	Over	Why
Focus areas separate from the risk register	One list for both	A scoring-derived "lowest items" list and a technical / relationship / execution risk register answer different questions
Customer view hides methodology	Showing weights, letters and notes	Northstar's VPs need what each pillar measures and its score; the acronym and weights are our method
Exact Fireworks design language	A generic dashboard theme	The audience is Fireworks; fidelity signals care. The customer view mirrors what Northstar would see from Fireworks in production
Data bundled at build time	Loading the CSV at runtime	Nothing can fail to fetch. A malformed CSV shows an error panel naming the problem; replacing the file and rebuilding updates both dashboards